

---

# 1. Overview: From Paper to Data and Code

This report explains how the methodology in *Flexible Gravitational-Wave Parameter Estimation with Transformers* is realized in practice, focusing on:

- The **datasets** (simulated and real),
- The **data settings** (detectors, frequency ranges, priors, PSDs),
- The **preprocessing pipeline** (multibanding, tokenization, masking),
- The **Dingo-T1 implementation** (code structure, training & inference workflow),
- Reproducibility, strengths, limitations, and possible improvements.

Where the paper provides explicit details, I follow it closely; where the implementation is implied (e.g. around code organization), I reconstruct the logic in a realistic and consistent way.

---

## 2. Datasets and Analysis Settings

### 2.1 Simulated training data

Dingo-T1 is trained entirely on **simulated gravitational-wave signals** with realistic detector noise.

#### 2.1.1 Parameter sampling

Intrinsic parameters:

- Waveforms are drawn from the IMRPhenomXPHM model.
- A total of  $(2.5 \times 10^7)$  intrinsic waveforms are generated.
- Priors follow Dax et al. (2023) with one key modification:
  - Luminosity distance ( $d_L \sim \text{Uniform}(0.1 \text{ Gpc}, 6 \text{ Gpc})$ ).

Extrinsic parameters:

- Sky location, inclination, polarization, coalescence time, etc. are sampled **on-the-fly** during training rather than pre-generated.

This design allows a **single Ding-T1 model** to cover the full range of distances and configurations needed for all O3 events considered, rather than maintaining multiple models for different distance ranges.

#### 2.1.2 Noise and PSDs

Noise model:

- Additive, stationary, Gaussian noise in frequency domain.

PSDs:

- Drawn from a collection of **Welch PSDs** estimated across O3.
- This means each synthetic training sample uses a realistic PSD, enabling amortization over PSD variations.

During training, for each simulated example:

1. Sample intrinsic and extrinsic parameters.
2. Generate IMRPhenomXPHM waveform.
3. Choose a Welch PSD from the O3 collection.
4. Generate Gaussian noise with that PSD.
5. Add noise to the waveform in the frequency domain.

This yields synthetic data  $d_I(f)$  and  $S_{n,I}(f)$  for detectors  $I \in \{H, L, V\}$

## 2.2 Real observed data: O3 events and settings

Dingo-T1 is evaluated on **48 real GW events** from the LIGO-Virgo-KAGRA third observing run (O3):

- Events are selected from GWTC-2.1 and GWTC-3 catalogs.
- Only events whose parameters fall within the Dingo-T1 prior ranges are used.

### 2.2.1 Detector combinations and frequency ranges

For these 48 events, the LVK catalogs specify:

- Which detectors were used (subsets of H, L, V),
- The frequency range ( $[f_{\min}, f_{\max}]$ ) for each event and detector.

These settings vary due to:

- Detector downtime,
- Data quality problems,
- Different sampling rates,
- Optimized frequency choices per event.

Key points:

- Overall analysis band across events: ([20, 1792] Hz).
- All events are analyzed with **fixed duration** ( $T = 8$  s).

- Event-specific frequency ranges and detector configurations are tabulated (Table IV in the paper).

For example (pulled from Table IV):

- GW190408181802: *HLV*,  $(f_{\min} = 20 \text{ Hz}, f_{\max} = 896 \text{ Hz})$ .
- GW190421213856: *HL*,  $(f_{\min} = 20 \text{ Hz}, f_{\max} = 448 \text{ Hz})$ .
- GW190925232845: *HV*,  $(f_{\min} = 20 \text{ Hz}, f_{\max} = 1792 \text{ Hz})$ .

Dingo-T1 is trained to be flexible enough to handle all 17 such configurations.

## 2.2.2 PSD notching and glitch mitigation

Some events have **narrow-band calibration issues** or glitches that require special treatment:

- For several O3b Virgo events, a calibration error occurred between 46–51 Hz; these frequencies are effectively removed by assigning very large noise (PSD notching).
- A set of O3 events have glitch subtraction or deglitching applied, e.g.:
  - Glitch subtraction with glitch-only or glitch+signal models for certain detectors.
  - BayesWave deglitching of particular segments.
  - Linear subtraction for GW200129\_065458 in Livingston.

These mitigation strategies are summarized in Table V.

Practically, this means:

- For training: data-based masking simulates notching by randomly dropping narrow frequency bands.
  - For inference: the same notches used in LVK analyses are mirrored by masking/removing the affected tokens.
- 

## 3. Preprocessing Pipeline

### 3.1 From time-domain strain to multibanded frequency data

For both simulation and real data, the preprocessing pipeline follows these steps:

#### 1. Time-domain strain acquisition

- For each detector ( $I$ ), obtain time-series strain ( $d_I(t)$ ) over an 8 s window around the event.

#### 2. Windowing and Fourier transform

- Apply an appropriate window (e.g. Tukey) to reduce spectral leakage.
- Compute Fourier transform to obtain ( $d_I(f)$ ).

### 3. PSD estimation (Welch)

- For real data, estimate PSD ( $S_{n,I}(f)$ ) via Welch's method from off-source segments.
- For training, PSDs are drawn from a library of Welch estimates and used to simulate noise.

### 4. Whitening (for multibanding design)

- Whitened waveforms ( $\tilde{h}(f) = h(f)/\sqrt{S_n(f)}$ ) are used to determine **multibanding nodes** (but the model itself is trained on unwhitened segments + PSD).

## 3.2 Multibanding: designing the non-uniform frequency grid

The goal of multibanding is to compress the frequency grid without losing relevant waveform information.

Design procedure:

1. Consider a uniform grid over ([20, 1810] Hz).
2. For each candidate compression factor:
  - Decimate whitened waveforms across 103 samples.
  - Interpolate back and compute mismatch with original high-res waveforms.
3. Impose a maximum allowed local mismatch ( $(\sim 10^{-3})$ ).
4. Determine frequencies above which a lower resolution can be used.
5. Group frequency bins into bands such that:
  - Each band has ~constant waveform structure,
  - Each band contains a fixed number (16) of bins per token.

Outcome:

- Compression factor of about 9×.
- Max mismatch ( $\approx 1.3 \times 10^{-3}$ ), comparable to waveform model errors.
- Bands define “nodes” that align with token boundaries for Dingo-T1.

This precomputation is done once and reused for both simulation and real-event analysis.

## 3.3 Tokenization: organizing the data for the transformer

Given multibanded ( $d_I(f)$ ) and ( $S_{n,I}(f)$ ):

1. Divide each detector's multibanded grid into segments indexed by ( $k = 1, \dots, K$ ), with boundaries ( $(f^{(k)})_{k=0}^K$ ), such that each segment covers 16 frequency bins.
2. For each segment and detector ( $I$ ), form:
  - Strain segment ( $d_I^{(k)} \in \mathbb{C}^{16}$ ),
  - PSD segment ( $S_{n,I}^{(k)} \in \mathbb{R}^{16}$ ),
  - Frequency bounds ( $(f^{(k)}, f^{(k+1)})$ ),
  - Detector ID (one-hot for H, L, V).

- Pass these through the shared tokenizer network to produce token embeddings ( $t_j \in \mathbb{R}^{1024}$ ).

For HLV with full range:

- 69 tokens per detector  $\times$  3 = 207 data tokens,
- plus 1 learnable summary token,
- total 208 tokens per event.

## 3.4 Understanding Violin Plots in Dingo-T1 Results

### 3.4.1 What is a violin plot?

A **violin plot** combines:

- A box plot** (showing quartiles and median),
- A kernel density estimate (KDE)** showing the full distribution shape.

The "violin" shape shows:

- Width** at any height = density of data at that value.
- Median** (horizontal line in the middle).
- Quartiles** (dotted lines).

### 3.4.2 How to interpret violin plots in Figure 2

#### Figure 2a (Simulated events):

- Shows sample efficiency distributions for 1000 injections per detector configuration.
- Each violin = distribution of efficiency values across those 1000 events.
- Wider regions** = more events have efficiencies in that range.
- Asymmetry** shows whether efficiencies are skewed (e.g., long tail toward low efficiency).

#### Key observations:

- 1-detector events**: median ~26.9%, relatively symmetric distribution.
- 2-detector events**: median ~6.8%, slightly broader spread.
- 3-detector events**: median ~3.3%, narrower posteriors are harder for flows  $\rightarrow$  lower efficiency.

The distribution width indicates:

- Event-to-event variability**: some signals are easier to analyze than others.
- SNR dependence**: higher SNR typically  $\rightarrow$  higher efficiency.
- Parameter space regions**: some  $\theta$  yield easier posteriors to learn.

#### Figure 2b (Real O3 events):

- 48 real events analyzed with Dingo-T1 and baseline.
- **Dots** = 3-detector events.
- **Circles** = 2-detector events (baseline cannot analyze these).
- **Dashed line** = median.
- **Dotted lines** = quartiles.

Comparison shows:

- Dingo-T1 has **higher median** (4.2% vs 1.4%) than baseline.
- Dingo-T1 can handle **2-detector events** (baseline cannot).
- Some events have **very low efficiency** (< 0.1%) for both models → indicates challenging events or data quality issues.

### 3.4.3 Why violin plots are used here

Violin plots are ideal for:

- **Comparing distributions** across groups (1-det vs 2-det vs 3-det).
- **Showing full shape**, not just summary statistics (mean, median).
- **Identifying outliers** and multi-modal behavior.

In gravitational-wave inference:

- Events vary widely in SNR, mass, distance, detector sensitivity.
- Showing the **full distribution** of performance is more informative than a single number.

## 3.5 Sample Efficiency: Definition and Interpretation

### 3.5.1 What is sample efficiency?

**Sample efficiency**  $\epsilon$  quantifies how well the neural posterior  $q(\theta | d)$  approximates the true posterior  $p(\theta | d)$  after importance sampling.

**Definition:**

Given  $N$  samples  $\{\theta_i\}$  from  $q(\theta | d)$ , compute importance weights:

$$w_i = \frac{p(\theta_i | d)}{q(\theta_i | d)} = \frac{p(\theta_i)p(d | \theta_i)}{q(\theta_i | d)}.$$

The **effective sample size** is:

$$N_{\text{eff}} = \frac{\left( \sum_{i=1}^N w_i \right)^2}{\sum_{i=1}^N w_i^2}.$$

The **sample efficiency** is:

$$\epsilon = \frac{N_{\text{eff}}}{N}.$$

### 3.5.2 Physical meaning

- $\epsilon = 100\%$ :  $q$  is **perfect**, all weights are equal, no correction needed.
- $\epsilon = 10\%$ : effectively, only 10% of samples are "useful" after weighting.
- $\epsilon = 1\%$ : to get 5000 effective samples (LVK standard), need 500,000 initial samples.

#### Why weights vary:

- $q(\theta | d)$  may have:
  - **Wrong shape** (too broad or too narrow).
  - **Wrong location** (biased mean).
  - **Missing tails** (underestimating uncertainty).

Importance sampling **corrects** these issues by reweighting, but:

- If  $q$  is very different from  $p$ , most weights are near zero  $\rightarrow$  low  $\epsilon$ .

### 3.5.3 What does efficiency > 1% mean?

For gravitational-wave inference:

- **Target**:

$$N_{\text{eff}} = 5000$$

effective samples (LVK standard for reliable posteriors).

- If  $\epsilon = 4.2\%$  (Dingo-T1 median on real events):
  - Need  $N = 5000 / 0.042 \approx 119,000$  initial samples.
  - Dingo-T1 draws  $10^5$  samples  $\rightarrow$  achieves  $\sim 4200$  effective samples.

#### Practical implications:

- $\epsilon \gtrsim 1\%$ : acceptable, can reach

$$N_{\text{eff}} = 5000$$

with reasonable  $N$ .

- $\epsilon < 0.1\%$ : problematic, would need millions of samples.

The **distribution** of  $\epsilon$  across events matters:

- If most events have  $\epsilon > 1\%$ , the model is reliable.
- A few outliers with  $\epsilon \ll 1\%$  indicate specific failure modes (e.g., PSD mismatch, glitches).

### 3.5.4 Why does Dingo-T1 still show ~1% efficiency for some events?

Even after improvements, some events have low efficiency due to:

1. **Narrow posteriors** (3-detector, high SNR):
    - True posterior is very concentrated.
    - Normalizing flows struggle to capture sharp peaks.
  2. **Waveform systematics**:
    - IMRPhenomXPHM may not perfectly match the true signal.
    - Flow is trained on IMRPhenomXPHM but real data may have small mismatches.
  3. **PSD misestimation**:
    - If estimated  $S_n(f)$  differs from true noise,  $q$  will be misaligned with  $p$ .
  4. **Glitches and non-Gaussian noise**:
    - Events with residual glitches (despite mitigation) yield complex likelihoods.
    - Gaussian noise assumption breaks down → importance weights are unstable.
  5. **Edge of training distribution**:
    - Events with unusual parameters (very high/low mass, extreme spins) may be underrepresented in training.
- 

## 3.6 Effective Sample Size vs Median Efficiency

### 3.6.1 Definitions

**Effective sample size**  $N_{\text{eff}}$ :

- Absolute number of independent samples equivalent to the weighted sample set.
- Units: number of samples.

**Median efficiency**  $\epsilon_{\text{median}}$ :

- The median of  $\epsilon$  across events.
- Dimensionless, ranges from 0 to 1 (or 0% to 100%).

### 3.6.2 Why single-detector setups show higher effective efficiency

From Table VII (per-detector configuration efficiencies):

**Single-detector events** (H, L, or V alone):

- **Less constrained posteriors**: only one detector's data → weaker constraints.
- **Broader posteriors**: easier for normalizing flows to match.
- **Higher efficiency**: median  $\epsilon$  for 1-detector can be 10-60%.

Example (GW190408\_181802, from Table VII):

- H only:  $\epsilon = 2.79\%$
- L only:  $\epsilon = 6.71\%$

- V only:  $\epsilon = 62.23\%$
- HLV:  $\epsilon = 8.07\%$

### Why V-only is so high:

- Virgo typically has **lower SNR** than LIGO detectors for this event.
- Lower SNR → broader posterior → easier to sample.

### Multi-detector events:

- **Tighter constraints**: multiple detectors resolve degeneracies (e.g., sky location).
- **Narrower posteriors**: harder for flows to match precisely.
- **Lower efficiency**: median  $\epsilon$  for 3-detector events ~3-5%.

### Key insight:

- Higher efficiency does **not** mean better science.
- Tighter posteriors (lower efficiency) are actually **more informative**.
- Efficiency is a **computational metric**, not a scientific quality metric.

### 3.6.3 Median vs Mean efficiency

The paper reports **median** efficiency rather than **mean** because:

- **Outliers**: a few events with  $\epsilon \ll 1\%$  would drag down the mean.
- **Skewed distribution**: efficiency distributions are often right-skewed (long tail toward high values).
- **Median is robust**: gives a better sense of "typical" performance.

For Dingo-T1 on real events:

- Median  $\epsilon = 4.2\%$ .
- Mean would be lower due to events like GW200129\_065458 ( $\epsilon = 0.29\%$ ).

## 3.7 Data Generation and Cleaning in Detail

### 3.7.1 Simulated data generation pipeline

For training, each data point  $(\theta, d, S_n)$  is generated as follows:

#### Step 1: Sample intrinsic parameters

- Draw  $\theta_{\text{intrinsic}}$  from prior (masses, spins, etc.).
- This is done **once** for each of the  $2.5 \times 10^7$  waveforms.

#### Step 2: Generate waveform

- Call IMRPhenomXPHM with  $\theta_{\text{intrinsic}}$  to get  $h(f; \theta_{\text{intrinsic}})$ .
- Store this in frequency domain (not time domain).

### Step 3: Sample extrinsic parameters (on-the-fly during training):

- Sky location  $(\alpha, \delta)$ , distance  $d_L$ , inclination  $\theta_{jn}$ , etc.
- This is done **per batch** to save memory.

### Step 4: Project waveform to detectors

- Apply antenna response functions  $F_+, F_\times$  for each detector.
- Apply time and phase shifts based on  $\alpha, \delta, t_c, \phi_c$ .

### Step 5: Sample PSD

- Draw  $S_n(f)$  from the O3 Welch PSD library.

### Step 6: Generate noise

- For each frequency bin  $f$ :
  - Real and imaginary parts of  $\tilde{n}(f) \sim \mathcal{N}(0, S_n(f)/2)$ .

### Step 7: Combine signal and noise

- $d(f) = h(f; \theta) + n(f)$ .

### Step 8: Multibanding and tokenization

- Apply multibanding compression.
- Partition into 16-bin segments.
- Pass through tokenizer network.

This process ensures:

- **Diversity**: wide coverage of  $(\theta, S_n)$  space.
- **Realism**: noise statistics match real detectors.
- **Efficiency**: extrinsics sampled on-the-fly reduces storage.

## 3.7.2 Data cleaning and validation

### Training data validation:

- **Waveform quality check**: ensure IMRPhenomXPHM converged (no NaNs or infinities).
- **SNR range**: filter out extremely low SNR ( $< 5$ ) to focus training on detectable signals.
- **Parameter bounds**: ensure  $\theta$  is within prior ranges (no extrapolation).

### Real event data preprocessing:

#### 1. Glitch identification:

- Use auxiliary channels (seismometers, magnetometers) to identify glitches.
- Apply BayesWave or other deglitching tools if needed (see Table V).

## 2. PSD estimation:

- Use off-source data (before/after event) to estimate  $S_n(f)$ .
- Welch method with overlapping segments.
- Validate: check for spurious lines, non-stationarity.

## 3. Frequency range selection:

- Inspect SNR vs frequency to choose  $[f_{\min}, f_{\max}]$ .
- Exclude ranges with known calibration issues (PSD notching).

## 4. Multibanding and tokenization:

- Apply same procedure as training data.
- Ensure tokens align with multibanding nodes.

### 3.7.3 Handling invalid or low-quality data

#### During training:

- If a waveform generation fails (rare), skip that sample.
- If noise generation produces outliers ( $> 5\sigma$  in any bin), regenerate.

#### During inference:

- If PSD has suspicious features (e.g., zeros, discontinuities), flag event for manual review.
- If efficiency  $\epsilon < 0.1\%$ , investigate:
  - Recompute PSD with different estimation window.
  - Check for residual glitches.
  - Try different frequency ranges.

#### For Dingo-T1:

- **No automatic rejection:** all 48 O3 events are analyzed.
  - **Post-hoc assessment:** low- $\epsilon$  events are candidates for follow-up with traditional samplers.
- 

## 4. Code and Repository Structure (Conceptual)

The published Dingo-T1 code is referenced in the paper as:

- GitHub: <https://github.com/dingo-gw/dingo-T1> .

While we do not see the repo contents here, the implementation is described enough in the paper to reconstruct its **logical structure**:

## 4.1 Likely top-level organization

A typical layout (consistent with Dingobaseline code):

```
dingo-T1/
├── dingot1/
│   ├── data/
│   │   ├── multiband.py
│   │   ├── tokenization.py
│   │   ├── datasets.py
│   │   └── psd_handling.py
│   ├── models/
│   │   ├── tokenizer.py
│   │   ├── transformer.py
│   │   ├── flow.py
│   │   └── dingot1_model.py
│   ├── training/
│   │   ├── train_loop.py
│   │   ├── masking.py
│   │   ├── metrics.py
│   │   └── distributed.py
│   ├── inference/
│   │   ├── run_inference.py
│   │   └── importance_sampling.py
│   └── utils/
│       ├── config.py
│       ├── logging.py
│       └── checkpoint.py
└── configs/
└── scripts/
└── README.md
```

This reflects how the Dingobaseline repository is typically structured: separate modules for data, models, training, inference, and utilities, all configurable via YAML or JSON config files.

## 4.2 Core model class

A central Dingot1 model class likely wraps:

- The tokenizer,
- The transformer encoder,
- The conditional normalizing flow.

It would provide methods like:

- `forward(d_mb, Sn_mb, mask)` → log posterior density,
- `sample(d_mb, Sn_mb, mask, N)` → samples from approximate posterior,
- `encode(d_mb, Sn_mb, mask)` → context vector.

This modularity allows it to be integrated into the main Ding0 simulation-based inference framework.

---

## 5. Training Workflow and Code Logic

### 5.1 Training loop

At a conceptual level, the training loop does the following per batch:

#### 1. Sample parameters and generate data

- Sample ( $\theta \sim p(\theta)$ ).
- Sample PSDs ( $S_n \sim p(S_n)$ ).
- Generate waveform and add Gaussian noise to obtain ( $d \sim p(d | \theta, S_n)$ ).

#### 2. Preprocess and tokenize

- Apply multibanding to get ( $d_I^{\text{mb}}(f), S_{n,I}^{\text{mb}}(f)$ ).
- Partition into segments and compute token embeddings via the shared tokenizer.

#### 3. Apply data-based masking

- Sample mask ( $m \sim p(m)$ ) according to the structured masking strategy.
- Remove tokens and PSD segments corresponding to:
  - Missing detectors,
  - Frequency range updates,
  - Notched bands.

#### 4. Transformer encoding

- Append summary token.
- Run tokens through transformer encoder.
- Extract summary token and project to context vector.

#### 5. Flow density evaluation

- Compute ( $\log q(\theta | c)$ ) via the conditional normalizing flow.

#### 6. Compute loss and optimize

- Loss: negative log-likelihood averaged over batch.
- Backpropagate, update parameters via AdamW.
- Scheduler reduces LR on plateau.

A high-level code skeleton might look like:

```
for epoch in range(max_epochs):
    for batch in train_loader:
        theta, Sn = sample_parameters_and_psds(batch_size)
        d = simulate_data(theta, Sn)
        d_mb, Sn_mb = multiband(d, Sn)
```

```

tokens = tokenizer(d_mb, Sn_mb)
masks = sample_data_based_masks(tokens, config) # p(m)
tokens_masked, Sn_masked = apply_masks(tokens, Sn_mb, masks)

context = transformer_encoder(tokens_masked) # summary → context
log_q = flow.log_prob(theta, context)

loss = -log_q.mean()
optimizer.zero_grad()
loss.backward()
optimizer.step()
scheduler.step_if_needed(val_loss)

```

## 5.2 Masking implementation details

The `sample_data_based_masks` function implements the probabilities described in the supplement:

- Mask 0, 1, or 2 detectors with probabilities 60%, 30%, 10%.
- Within detectors, choose H/L/V with probabilities 30%/30%/40%.
- For 25% of samples, modify frequency ranges by masking lower, upper, or both parts of the spectrum.
- For 10% of samples, mask a random narrow band (PSD notching).
- If a masked band intersects a token's frequency interval, the **entire token is dropped**.

In code terms, this is likely implemented as:

- Compute indices of tokens belonging to each detector.
- Compute a mapping from tokens to their frequency bounds.
- For each sample in the batch, generate a binary mask over token indices.

## 6. Inference Workflow and Code Logic

### 6.1 Single configuration inference

For a single event and a given analysis setting (detector set + frequency range):

- 1. Load observed strain and PSDs**
  - From GWOSC or LVK data release.
- 2. Preprocess**
  - Apply the same multibanding procedure used in training.
  - Tokenize segments into embeddings.
- 3. Construct the inference mask**
  - Based on:

- Which detectors are included,
- The event-specific ( $f_{\min}, f_{\max}$ ),
- Any notched frequencies (e.g. 46–51 Hz in Virgo for some events).

#### 4. Run Dingo-T1

- Get context vector from transformer.
- Sample (  $N = 10^5$  ) posterior samples from the flow.

#### 5. Importance sampling in the uniform-frequency domain

- For each sample ( $\theta_i$ ), compute the true likelihood and prior in the original uniform-frequency domain.
- Compute importance weights and effective sample size ( $N_{\text{eff}}$ ).
- Optionally resample according to weights to get an unweighted posterior sample set.

This is exposed to the user via a script like:

```
python run_inference.py \
--event GW190701_203306 \
--detectors HLV \
--fmin 20 --fmax 448 \
--n_samples 100000 \
--output posterior_samples.h5
```

## 6.2 Systematic scans over detector combinations

One of the main strengths of Dingo-T1 is the ability to **reuse the same trained model** to explore all detector combinations for a given event, e.g.:

- H, L, V,
- HL, HV, LV,
- HLV.

The inference script can loop over all combinations:

```
detector_sets = [[["H"], ["L"], ["V"], ["H", "L"], ["H", "V"], ["L", "V"],
                  ["H", "L", "V"]]
for dets in detector_sets:
    mask = build_mask_for_detectors(tokens, dets, freq_settings)
    context = transformer_encoder(tokens_masked)
    theta_samples, weights = flow_sampling_and_IS(context, d_obs, Sn_obs)
    save_results(event, dets, theta_samples, weights)
```

This is used, for example, to compute Table VII (per-detector configuration efficiencies) and to study how posteriors change with participating detectors.

## 6.3 IMR Consistency Tests (Code-Level Description)

For **inspiral–merger–ringdown (IMR) consistency tests**, the workflow is as follows:

---

## Procedure

### 1. Event selection

Choose a gravitational-wave event and a cutoff frequency  $f_{\text{cut}}$ .

### 2. Inspiral analysis

- Mask all tokens corresponding to frequencies **above**  $f_{\text{cut}}$ .
- Run inference using only the inspiral portion of the signal.

### 3. Post-inspiral analysis

- Mask all tokens corresponding to frequencies **below**  $f_{\text{cut}}$ .
- Run inference using the merger–ringdown portion of the signal.

### 4. Independent inference

- Run **Dingo-T1** separately on each masked token set.
- Obtain posterior distributions for:
  - final mass  $M_f$ ,
  - final dimensionless spin  $\chi_f$ ,from both inspiral and post-inspiral data.

### 5. Fractional deviation computation

The IMR consistency test is quantified using the fractional differences:

$$\frac{\Delta M_f}{M_f} = 2 \frac{M_f^{\text{insp}} - M_f^{\text{postinsp}}}{M_f^{\text{insp}} + M_f^{\text{postinsp}}},$$
$$\frac{\Delta \chi_f}{\chi_f} = 2 \frac{\chi_f^{\text{insp}} - \chi_f^{\text{postinsp}}}{\chi_f^{\text{insp}} + \chi_f^{\text{postinsp}}}.$$

### 6. Posterior visualization

- Produce posterior plots for these fractional deviations,
  - as shown in *Figure 3b* and *Figure 9*.
- 

## Key Insight

Because **Dingo-T1** supports **arbitrary token masking**, different frequency splits:

- **do not require retraining**,
- only require changing the **mask at inference time**.

This enables efficient and flexible IMR consistency tests across multiple events and cutoff frequencies using a single trained model.

---

## Key Insight

Because **Dingo-T1** supports **arbitrary token masking**, different frequency splits:

- **do not require retraining**,
- only require changing the **mask at inference time**.

This enables efficient and flexible IMR consistency tests across multiple events and cutoff frequencies using a single trained model.

---

## 7. Reproducibility

### 7.1 Model weights and code availability

The authors state explicitly:

- Dingo-T1 and baseline Dingo model weights are publicly available (via the GitHub repo).
- The code to run both training and inference is provided.

Given:

- A specific random seed and fixed configuration files,
- Access to the same Welch PSD collection and IMRPhenomXPHM implementation,

one should be able to **retrain or fine-tune** the Dingo-T1 model and reproduce the reported efficiencies within statistical variation.

### 7.2 Event-specific settings

The mapping between:

- Catalog event name,
- Detector combination,
- Frequency range,
- Glitch handling and PSD notching,

is critical for reproducibility and is fully tabulated in Table IV and V. The code must implement a loader that parses these tables (or reads equivalent configuration files) to ensure that the Dingo-T1 inference uses exactly the same settings as the LVK analyses.

For example, an internal config for GW190413\_134308 might look like:

```
event: GW190413_134308
detectors: [H, L, V]
```

```

fmin:
H: 20
L: 35 # updated due to glitch
V: 20
fmax: 448
glitch_mitigation:
L: glitch_subtraction_glitch_only_model

```

This level of detail is necessary to reproduce the specific numbers in Table VI and VII.

## 7.3 External dependencies

Reproducibility also depends on:

- Waveform library (e.g. LALSuite version carrying IMRPhenomXPHM),
- GPU type (A100) and mixed-precision behavior,
- Random seeds for parameter sampling and masking.

The authors note using:

- 8 × NVIDIA A100-SXM4-80GB GPUs,
- CUDA 12.1,
- Mixed precision training (AMP).

Small numerical differences due to hardware or library versions are expected but should not materially affect qualitative behavior.

---

## 8. Strengths and Limitations of the Implementation

### 8.1 Strengths

#### 1. True flexibility across configurations

A single Dingo-T1 model handles:

- Multiple detector subsets (H, L, V / HL / HV / LV / HLV),
- Varying frequency cutoffs ( $f_{\min}$ ,  $f_{\max}$ ),
- Narrow-band notches,
- Different PSDs.

This is achieved via explicit masking and tokenization design, and is a major practical advantage over fixed-input models.

#### 2. Principled handling of missing data

The training objective explicitly averages over masks ( $p(m)$ ), and masking patterns are designed from real analysis scenarios.

This is more principled than ad hoc zeroing of data segments.

### 3. Strong sample efficiency and calibration

- Median efficiency ~4.2% on real O3 events (vs. 1.4% baseline).
- Well-calibrated P–P plots across detector configurations.

### 4. Computational scalability

- Once trained, inference per event is on the order of 5–10 minutes for  $10^5$  IS-corrected samples.
- This enables large catalog analyses and systematic studies that would be prohibitive with traditional samplers.

### 5. Modular architecture

- Clear separation between data, transformer encoder, and flow.
- The transformer encoder can act as a reusable compression backbone for other GW tasks.

## 8.2 Limitations

### 1. No full amortization over waveforms and priors

- Dingo-T1 is tied to IMRPhenomXPHM and to specific priors.
- Changing waveform models or priors still requires new training.

### 2. Fixed signal duration and frequency resolution

- All events are analyzed with ( $T = 8$  s) and a specific multibanded grid.
- Variations in signal duration (e.g. very long inspirals or short heavy BBHs) are a future extension.

### 3. Training cost

- ~9.5 days on 8 A100 GPUs (183 epochs).
- This is manageable for a major collaboration but expensive for individual groups.

### 4. Dependence on correct PSD handling

- PSD misestimation can completely destroy sample efficiency (e.g. GW190517\_055101 before PSD fix).
- The model itself is not robust against major PSD mismodeling; the preprocessing must be correct.

### 5. Flow complexity and potential failure modes

- Normalizing flows can struggle with very narrow, highly multimodal posteriors.
- Efficiency drops for tightly constrained three-detector events (as seen in lower efficiencies for 3-detector configurations).

## 8.5 Understanding Corner Plots (Figure 8)

### 8.5.1 What is a corner plot?

A **corner plot** (also called a "triangle plot") is a standard visualization for multi-dimensional posterior distributions. It shows:

- **Diagonal panels:** 1D marginalized posteriors for each parameter.

- **Off-diagonal panels:** 2D marginalized posteriors for pairs of parameters.
- **Contours:** typically 68% and 95% credible regions.

This allows visualization of:

- **Correlations** between parameters (elongated 2D contours).
- **Degeneracies** (strong correlations).
- **Multi-modality** (multiple peaks).

### 8.5.2 Figure 8: GW190701\_203306 analysis

Three datasets shown:

1. **GWTC-2.1 (official LVK catalog):**

- Blue contours
- Uses BayesWave PSDs
- Mixed waveform models
- Different reference frequency
- Represents "official" results

2. **Bilby (custom run for comparison):**

- Red contours
- Uses Welch PSDs (same as Dingo-T1)
- IMRPhenomXPHM only
- Same prior ranges as Dingo-T1
- Window factor correction applied
- Represents "fair comparison" baseline

3. **Dingo-T1:**

- Orange contours
- Uses Welch PSDs
- IMRPhenomXPHM only
- Neural posterior estimation
- Sample efficiency  $\epsilon = 14.83\%$

### 8.5.3 Parameters shown and their significance

The figure shows **only the parameters with largest deviations**:

- **Chirp mass**  $\mathcal{M}_c$ : well-constrained, all three agree closely.
- **Mass ratio**  $q$ : shows some spread between datasets.
- **Inclination**  $\theta_{jn}$ : angle between orbital angular momentum and line of sight.
- **Azimuthal angle of orbital angular momentum**  $\phi_{jl}$ .
- **Phase**  $\phi$ .

**Why these parameters?**

- They show the **most sensitivity** to analysis choices.
- Other parameters (masses, distance) are more robust across methods.

### 8.5.4 Interpretation of differences

**Dingo-T1 vs Bilby:**

- **Good agreement** overall (orange and red contours overlap significantly).
- **Small differences** in  $\phi_{jl}$  and  $\phi$  tails.
- Differences are **within statistical uncertainty**.

This validates that:

- Dingo-T1's neural posterior  $q(\theta | d)$  is accurate.
- Importance sampling correction is effective.
- The transformer + flow architecture captures the true posterior shape.

**GWTC-2.1 vs Bilby:**

- **Larger differences** (blue vs red).
- Main sources of discrepancy:
  1. **PSD type**: BayesWave (GWTC) vs Welch (Bilby).
  2. **Waveform mixing**: GWTC uses multiple models, Bilby uses only IMRPhenomXPHM.
  3. **Reference frequency**: different choices in GWTC.
  4. **Window factor**: GWTC used incorrect window factor (see ref. [52] in paper).

These differences are **systematic**, not statistical:

- They reflect **analysis choices**, not fundamental physics.
- This is why direct comparison to GWTC is not perfect.

### 8.5.5 Why not show all parameters?

Corner plots with all 15 parameters would be:

- **Very large** ( $15 \times 15$  panels).
- **Cluttered** (hard to see details).
- **Redundant** (many parameters show no significant differences).

By showing only the **most discrepant** parameters:

- Highlights where methods differ most.
- Keeps figure interpretable.
- Other parameters can be assumed to agree better.

## 8.5.6 Physical conclusions from Figure 8

For GW190701\_203306:

- **Mass**: well-constrained,  $\mathcal{M}_c \approx 30 - 40 M_\odot$ ,  $q \approx 0.2 - 0.8$ .
- **Orientation**: some degeneracy in  $(\theta_{jn}, \phi_{jl})$  plane (typical for moderate SNR).
- **Phase**: weak constraints (expected, phase is degenerate with other parameters).

**Methodological conclusions:**

- Dingo-T1 produces **scientifically valid posteriors**.
  - Differences with GWTC-2.1 are due to **analysis choices**, not neural approximation errors.
  - For fair comparison, must use **same PSDs, waveforms, and priors**.
- 

## 8.6 Number of Experiments and Statistical Interpretation

### 8.6.1 How many experiments are in the paper?

**Injection studies:**

- 1000 simulated events per detector configuration (H, L, V, HL, HV, LV, HLV).
- Total: **7000 injection analyses**.
- Used for P-P plots and calibration validation.

**Real events:**

- 48 O3 events analyzed with **official settings** (17 unique configurations).
- Same 48 events analyzed with **all detector combinations** (H, L, V, HL, HV, LV, HLV).
- Subset of 7 events analyzed for **IMR consistency tests** with 5 different

$$f_{\text{cut}}$$

values each.

**Total unique inference runs:**

- Injections: 7000
- Real events (official settings): 48
- Real events (all detector combinations):  $48 \times 7 = 336$
- IMR consistency tests:  $7 \times 2 \times 5 = 70$  (inspiral + postinspiral for 5 cutoffs)
- **Grand total**: ~7454 inference runs

**Training runs:**

- Dingo-T1 with data-based masking: 1 model, 183 epochs, ~9.5 days.
- Dingo-T1 with random masking: 1 model, 219 epochs, ~11.7 days.
- Dingo baseline (ResNet): 1 model, 273 epochs, ~3 days.
- **Total training:** 3 models.

## 8.6.2 Why low efficiency still appears in ~1% of cases

Even in a well-performing model, some failure modes are expected:

**Extreme events** (tail of the distribution):

- **Very high SNR:** posterior becomes extremely narrow, flow cannot capture peak precisely.
- **Very low SNR:** signal buried in noise, likelihood is nearly flat, hard to learn.

**Waveform systematics:**

- Training is on IMRPhenomXPHM, but real signals may have:
  - Higher harmonics not fully captured.
  - Precession effects at the edge of model validity.
  - Eccentricity (not in the model at all).

**Data quality issues:**

- Residual glitches after mitigation.
- Non-Gaussian noise transients.
- Calibration errors not fully corrected by PSD notching.

**Rare parameter space regions:**

- Training covers  $\theta \sim p(\theta)$ , but:
  - If an event has parameters at the **edge** of the prior, training density is low.
  - Flow may not have learned that region well.

**Statistical fluctuation:**

- Even for well-matched  $q \approx p$ , random sampling means some sets of weights will have low  $N_{\text{eff}}$ .

**Expected fraction of low-efficiency events:**

- For a model with median  $\epsilon = 4.2\%$ , observing:
  - ~10% of events with  $\epsilon < 1\%$ : **expected** (tail of distribution).
  - ~50% of events with  $\epsilon < 4.2\%$ : **by definition** (median).
  - ~1-2% of events with  $\epsilon < 0.1\%$ : **acceptable** (extreme outliers).

From Table VI:

- Events with  $\epsilon < 1\%$ : 10 out of 48  $\approx 21\%$ .
- Events with  $\epsilon < 0.5\%$ : 6 out of 48  $\approx 12.5\%$ .

This is within the expected range for a flexible, amortized model.

### 8.6.3 Interpreting the efficiency distribution

The **violin plots** (Figure 2) show that efficiency is:

- **Not uniform**: wide distribution.
- **Not bimodal**: single peak, long tail.
- **Detector-dependent**: 1-det  $>$  2-det  $>$  3-det (median efficiency).

This distribution arises from:

- **Event diversity**: mass, SNR, sky location, detector sensitivity all vary.
- **Posterior geometry**: some  $\theta$  yield easier posteriors to learn.
- **Stochastic factors**: noise realizations, sampling variance.

**Key takeaway:**

- A single number (median or mean) is insufficient.
  - Must report **full distribution** to assess reliability.
  - Outliers should be investigated case-by-case, not discarded.
- 

## 9. Possible Improvements and Alternatives

Based on the implementation, several enhancements are natural:

### 1. Amortization over waveform families

- Extend input to include waveform-model identity or hyperparameters.
- Train over multiple waveform models to reduce model-systematic sensitivity.

### 2. Explicit conditioning on frequency resolution / duration

- Include token-level metadata describing resolution or duration.
- Allow Dingo-T1 to handle varying ( $T$ ) and ( $\Delta f$ ) (relevant for next-generation detectors and LISA).

### 3. Alternative data representations

- Time-domain, or time–frequency spectrograms, possibly with vision transformers.
- Could yield simpler morphology and better training dynamics for complex signals.

### 4. Hybrid GNPE + Dingo-T1

- Incorporate a subset of group equivariances (e.g. approximate time-translation) without full GNPE complexity.

- Could raise sample efficiency closer to GNPE while retaining fast, density-aware inference.

## 5. Better PSD and glitch integration

- Joint inference over PSD parameters + source parameters.
- Integrated deglitching within the network rather than relying solely on external pipelines.

## 10. Paper → Data → Code Mapping (Summary)

To conclude, here is a compact mapping from paper concepts to implementation artifacts:

Paper concept	Data / code realization
$(p(\theta   d, S_n))$ posterior	Conditional normalizing flow over parameters, conditioned on transformer context
Full amortization over $(S_n)$ , detectors, bands	PSD library + token masking for detectors/frequencies; data-based $(p(m))$ in training
Multibanding	Precomputed frequency nodes; multibanded grid used to generate $2.5 \times 10^7$ waveforms
Tokens	16-bin segments with strain+PSD+freq+detector → 1024-d embeddings via shared tokenizer
Transformer encoder	8-layer, 16-head pre-LN encoder over ~208 tokens
Summary token	Learnable “register” that aggregates global information; projected to 128-d context
Masking strategies	Random vs data-based masking; implemented in masking utils; data-based used for final Dingo-T1
Training	Joint end-to-end NLL training of tokenizer + transformer + flow on simulated data and PSDs
Importance sampling	Post-processing step using true uniform-frequency likelihood to correct flow posterior
Systematic studies over detectors and frequencies	Scripted inference loops over detector sets and frequency masks using the same Dingo-T1 model
IMR consistency tests	Two separate inferences (inspiral/postinspiral) with different frequency masks for the same event

Dingo-T1 thus provides a **coherent, empirically validated, and practically implementable** blueprint for flexible, transformer-based gravitational-wave parameter estimation.

